


UDP & TCP



Fundamentos de Redes

**Mestrado Integrado em
Engenharia de Computadores e Telemática
DETI-UA**

TCP/IP Reference Models



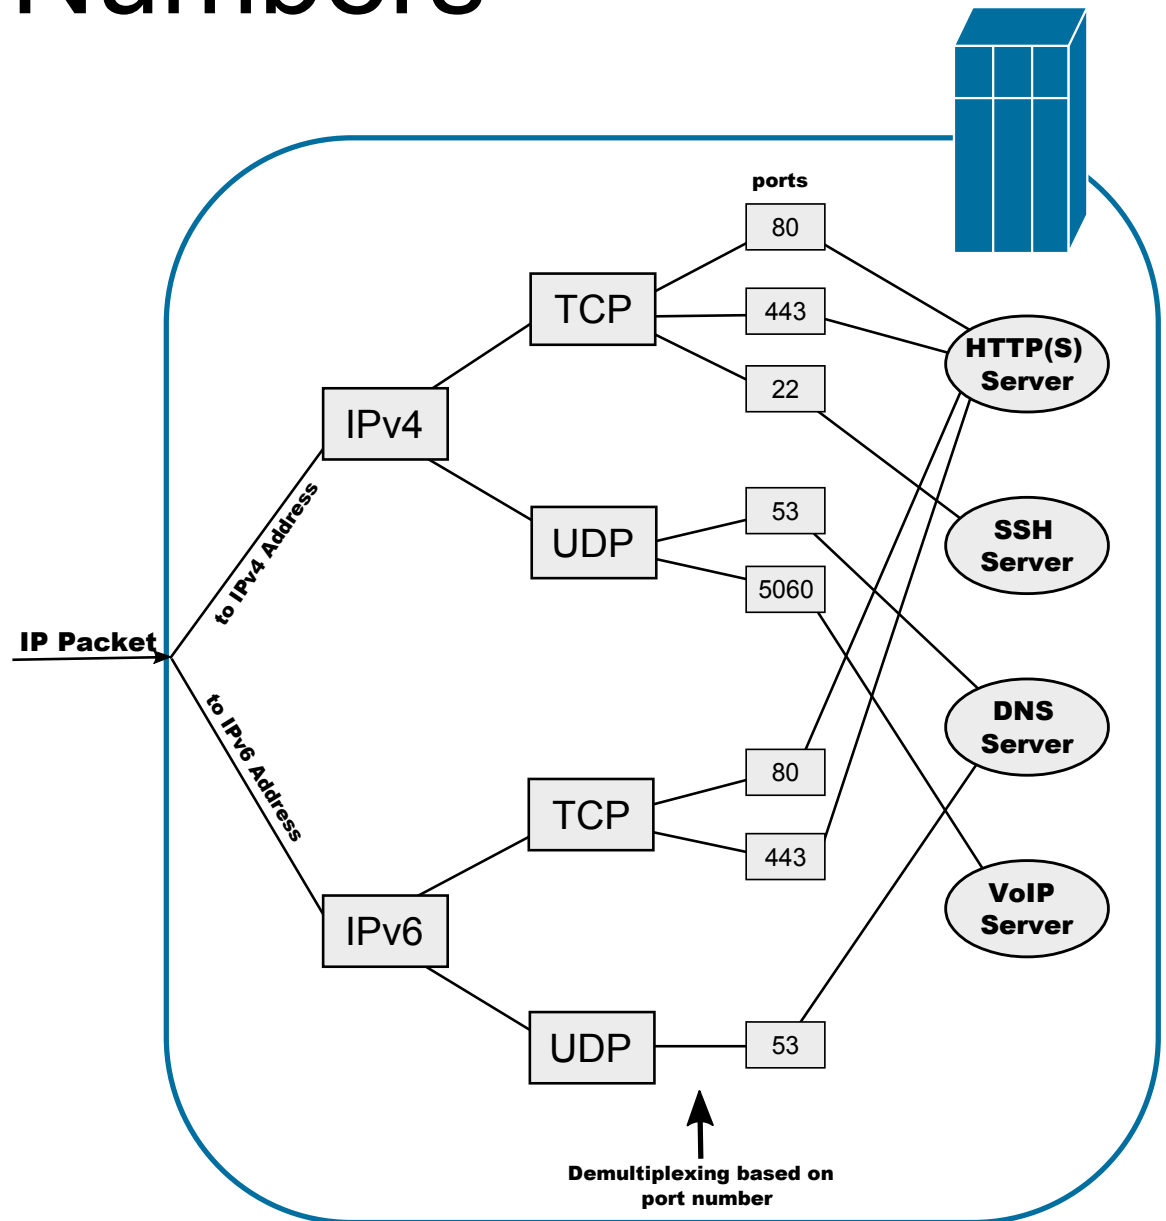
TCP/IP



OSI

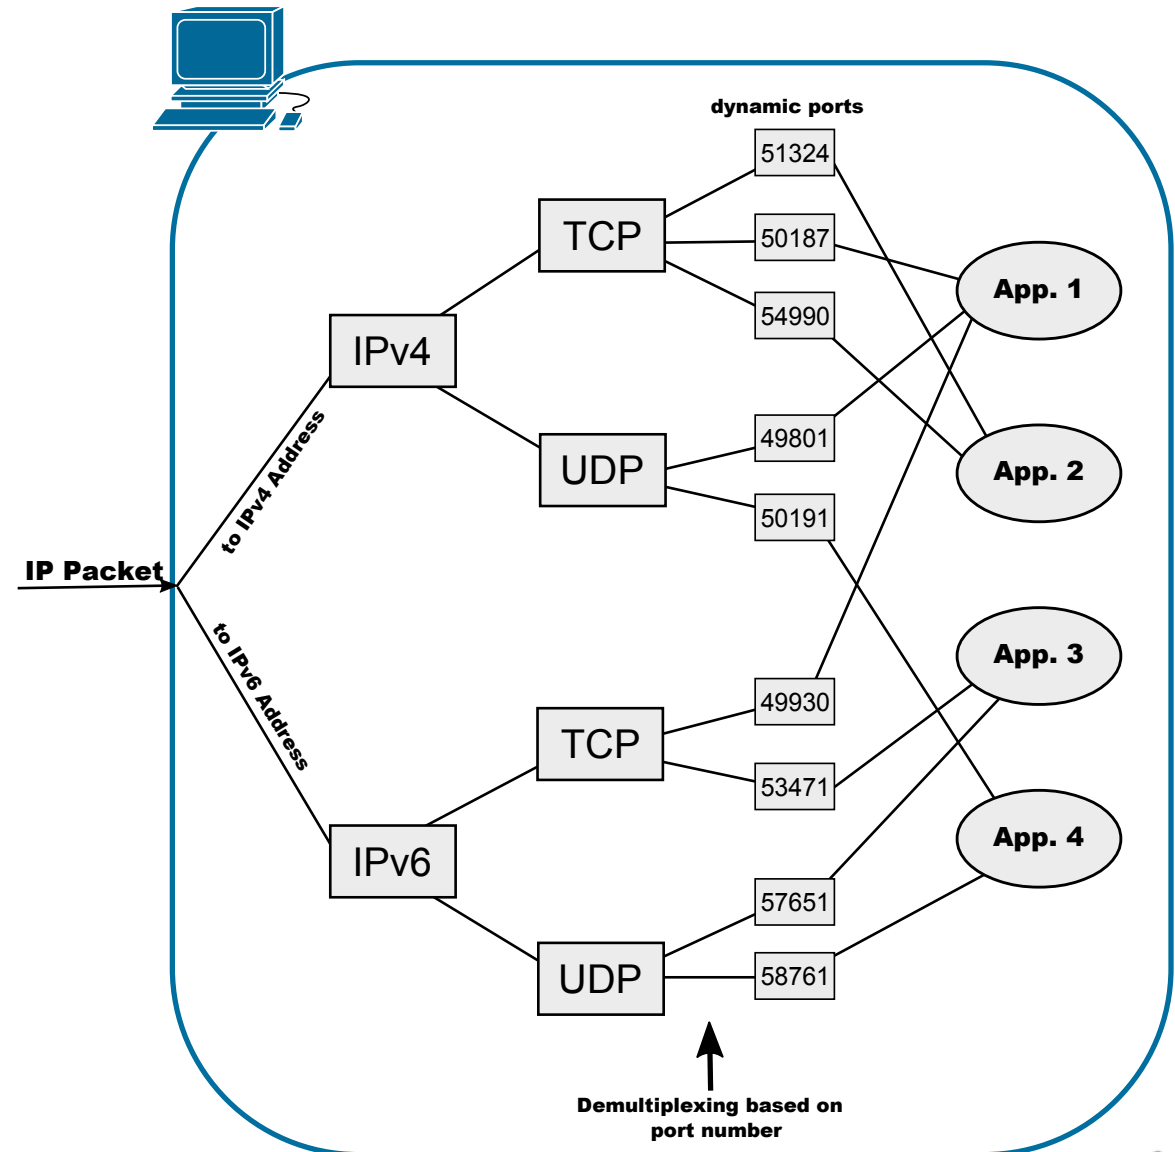
Port Numbers

- The packet destination IP address identifies a target host on the network.
- The transport protocol identifies the type of communication.
- The port number identifies the application running in the target host.
 - Chosen by the application/service.
 - ➔ Non-dynamic ports.
 - Each application/service may use more than one port.
 - The OS assures the assignment of different port numbers to each application.



Ephemeral/Dynamic Ports

- Ephemeral/Dynamic ports are used to identify a client application in a client-server communication.
- The Internet Assigned Numbers Authority (IANA) suggests the range 49152 to 65535.
- Randomly assigned by OS.



Well Known Port Numbers

Decimal Keyword		Protocol	Description
20	FTP-DATA	TCP	File Transfer Protocol (dados)
21	FTP-CONTROL	TCP	File Transfer Protocol (controlo)
22	SSH	TCP	Secure Shell (SSH) service
25	SMTP	TCP	Simple Mail Transport Protocol
67,68	BOOTP	UDP	Bootstrap Protocol (DHCP)
53	DNS	UDP/TCP	Domain Name System
69	TFTP	UDP	Trivial File Transfer Protocol
80	HTTP	TCP	Hypertext Transfer Protocol

- Many Internet IP services were the subject of study by IETF that proposed adequate support protocols.
- For these services, IETF together with the protocol specification proposed also a number (or numbers) to be used by that service at the server side.
- E.g., for protocol HTTP IETF recommends the usage of port 80. Therefore, all Web Browsers use port 80 as default for HTTP accesses.

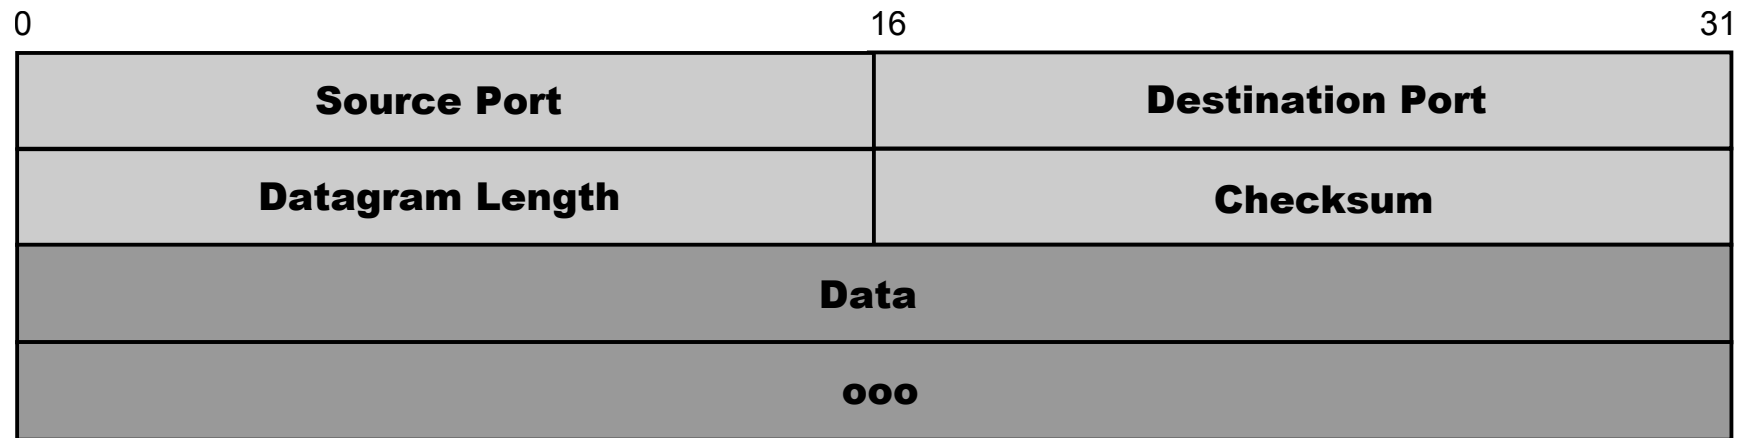


User Datagram Protocol (UDP)

- Provides a data transport service with the performance characteristics offered by the IP network.
- Provides exchange of data between individual applications, and not only between hosts, with the introduction of a port identifier field.
- Does not provide any mechanism to recover from lost messages.
- Does not provide any mechanism to order received data.
 - ◆ Must rely on information at the application level.
- Allows to send data to multiple destinations simultaneously (point-to-multipoint communications).



UDP Datagram (Header+Data)

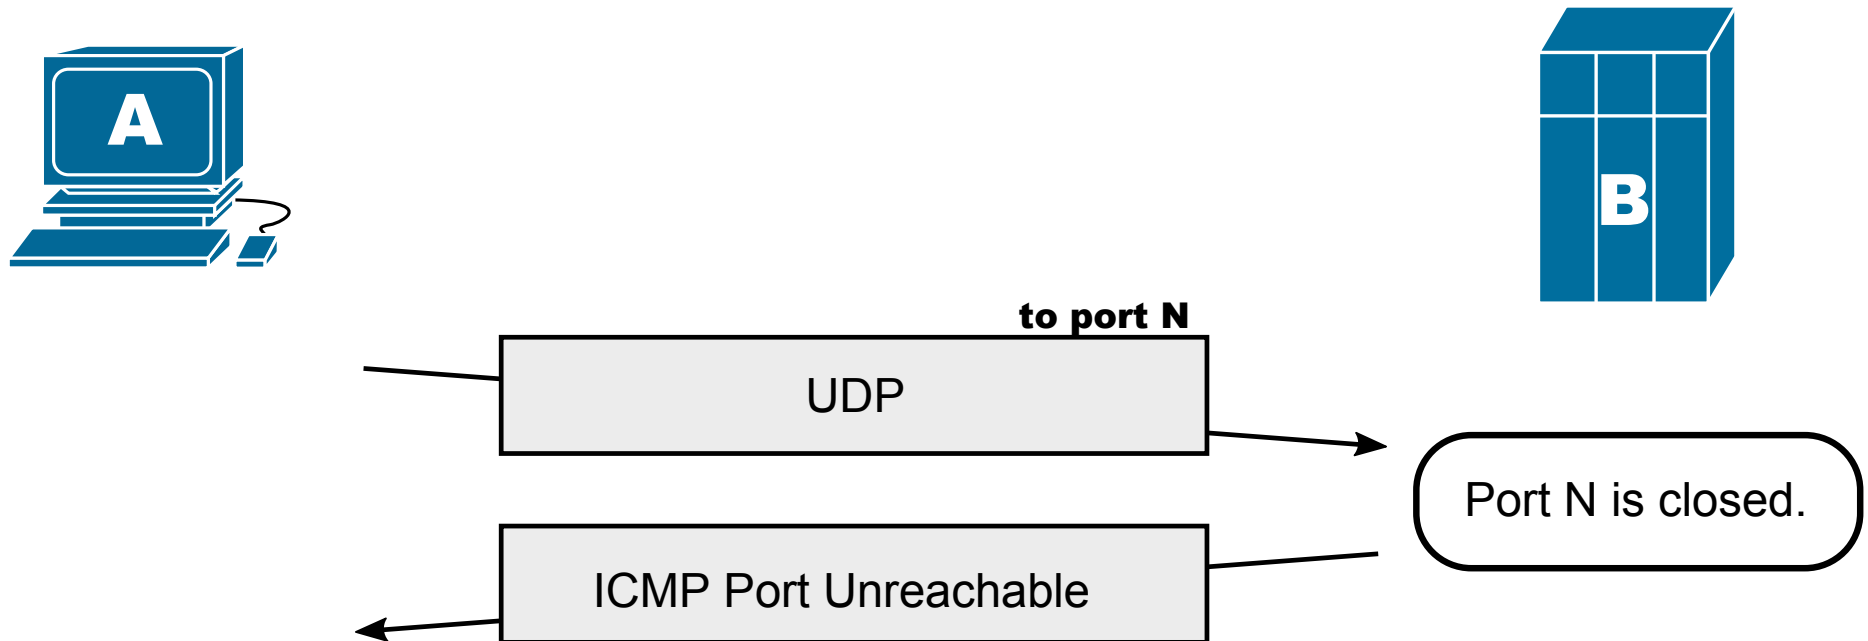


- Source Port (2 bytes): defines the port number assign/chosen by the sender application. This field is optional, if not used should be field with zeros.
- Destination Port (2 bytes): defines the port number assign/chosen by the receiver application.
- Datagram Length (2 bytes): defines the size, in bytes, of the datagram (header+data).
- Checksum (2 bytes): used for data error detection and validation at the end-points. This field is optional, if not used should be field with zeros.
 - ♦ The checksum is calculated based on the UDP datagram and a pseudoheader IP (IP protocol identifier, source and destination IP addresses, and length of the IP datagram).
 - ♦ Can be used to verify if the end-points are the correct ones.



UDP Closed Port

- When an UDP packet arrives to a host, but the UDP port is not open (no application listening):
 - ♦ The host responds with a packet ICMP *port unreachable*.

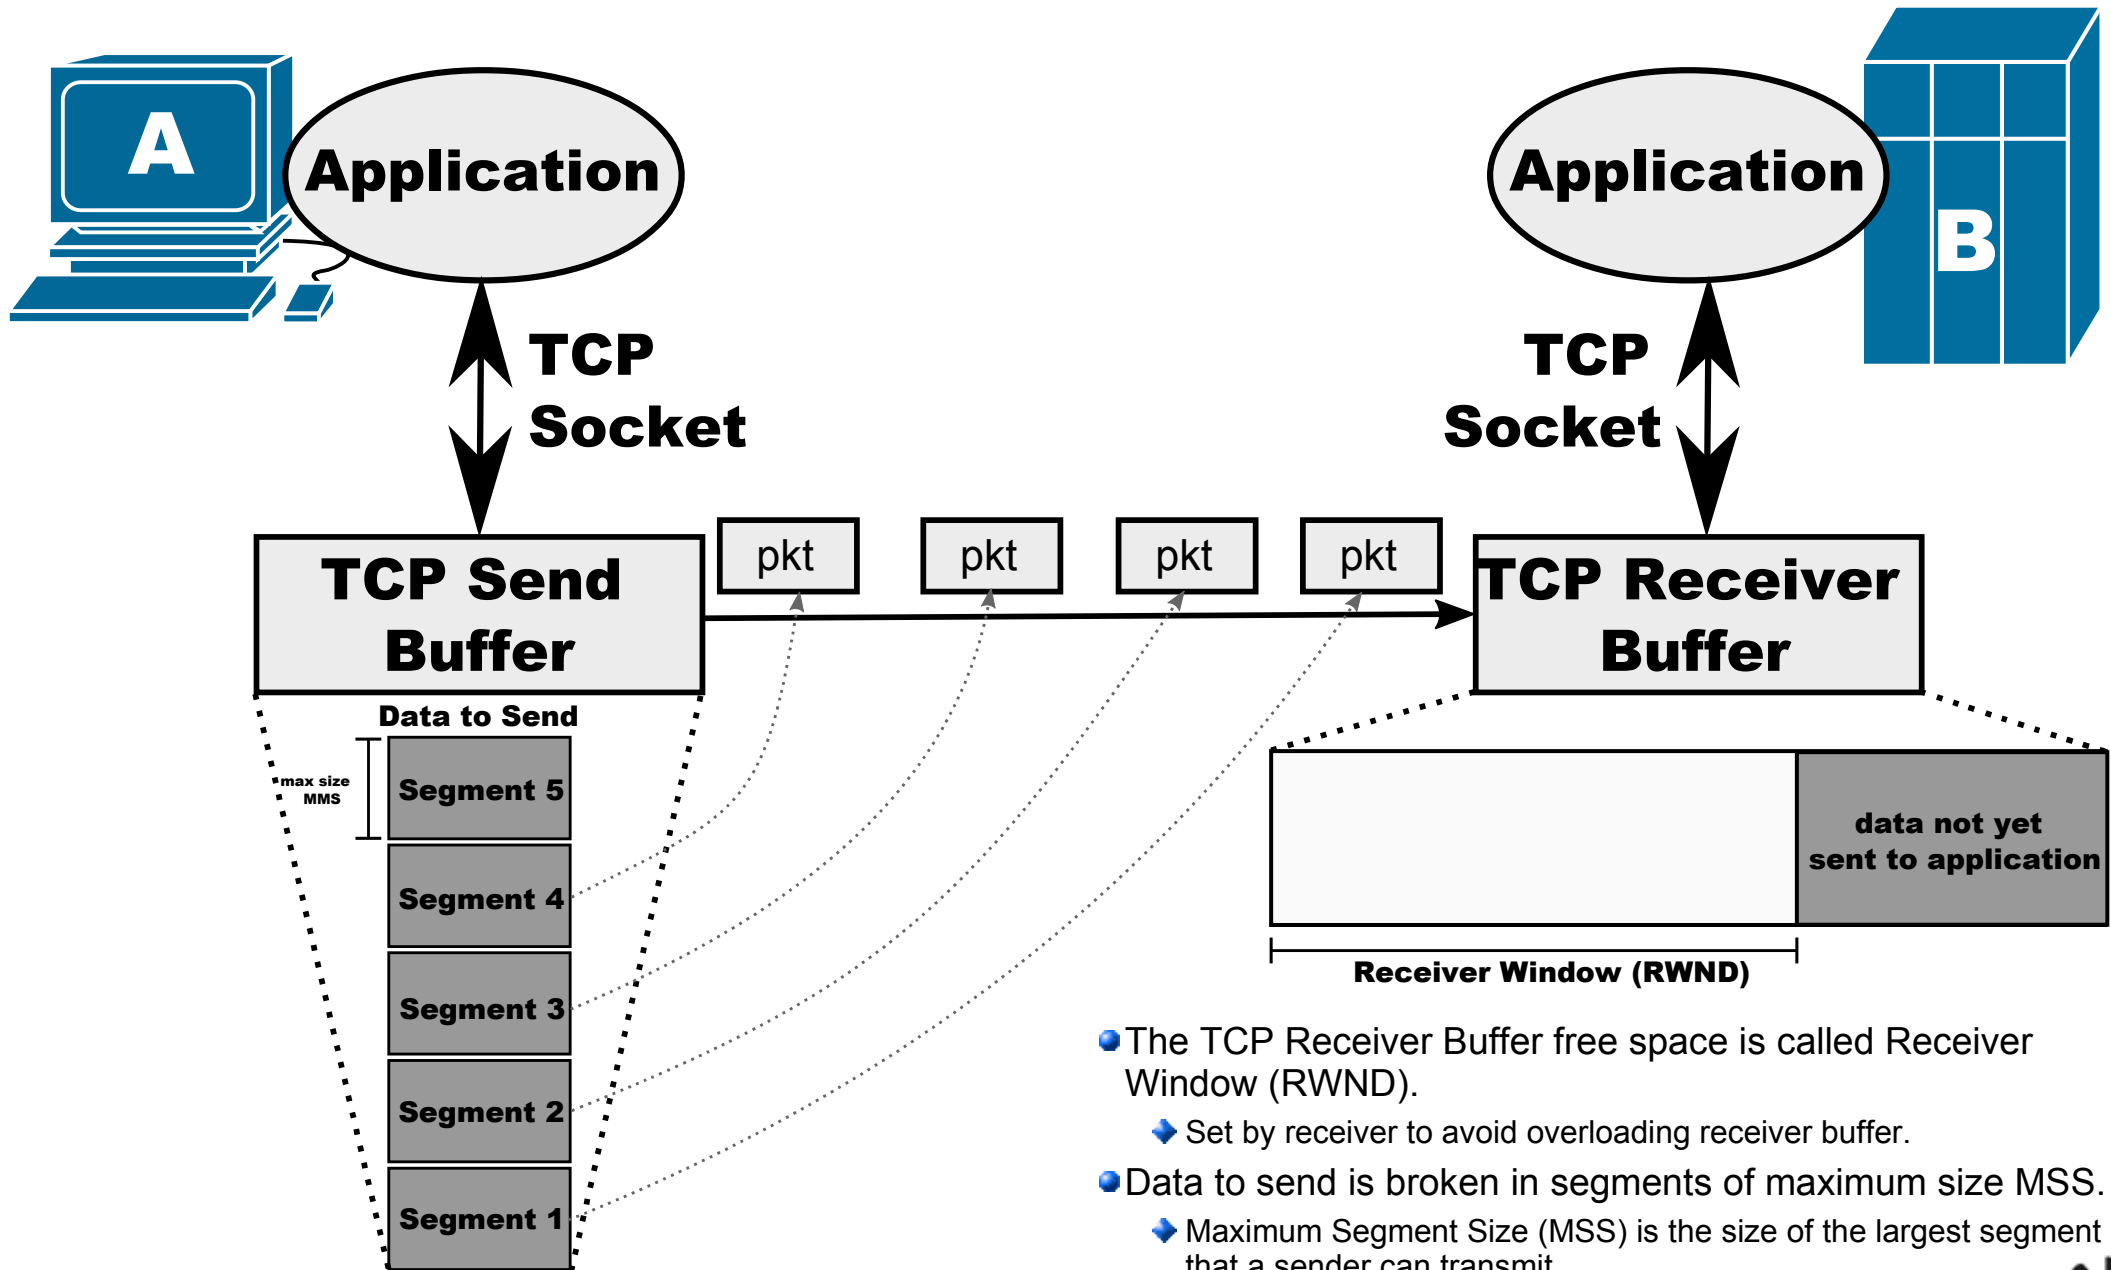


Transmission Control Protocol (TCP)

- Provides a reliable data transport service.
 - ♦ Data is received by the destination application without any losses and in order.
- Is connection-oriented protocol.
 - ♦ End-points establish a logical channel, to which are assigned applications' identifiers and the memory resources required to have a reliable transmission of data.
 - ♦ This connection is called Session.
- It is bi-directional.
 - ♦ Both end-points can send and receive data using the same logical channel.
- Traditional TCP supports only point-to-point connections.
 - ♦ See: Multipath TCP on last slide.
- Provides mechanisms to establish and terminate the connection.
- Network congestion and/or temporary lack of connectivity result in variable delays and consequent losses of packets (at transit or by timeout).
 - ♦ TCP includes algorithm that allows to efficiently react in this scenario.




TCP Buffers and Receiver Window



- The TCP Receiver Buffer free space is called Receiver Window (RWND).
 - ◆ Set by receiver to avoid overloading receiver buffer.
- Data to send is broken in segments of maximum size MSS.
 - ◆ Maximum Segment Size (MSS) is the size of the largest segment that a sender can transmit.
 - ◆ Defined by local configuration.



TCP Packet Header (1)



TCP Packet Header (2)


- The TCP Header has a variable length.
 - ♦ The field *Options* may contain additional fields.
- *Source Port* and *Destination Port* fields define the respective end-point ports.
- The *Sequence Number* defines the index of the first data byte in this segment.
- The *Acknowledge Number* defines the value of the next sequence number.
 - ♦ Acknowledges the good reception of data until byte with index (Acknowledge Number-1)
- *HLen* defines the size of the TCP header in bytes.
 - ♦ When *Options* are present, the field *Padding* is used to extend the header size to a multiple of 32 bytes.
- *Window* field defined the number of data bytes the sender of this segment is willing to accept (Receiver Window – RWND).
- *Control Bits* field is a binary set of flags
 - ♦ URG: Urgent Pointer is a valid field.
 - ♦ ACK: Acknowledgement is a valid field.
 - ♦ PSH: Data requires Push (receiver should push immediately data from TCP buffer to application).
 - ♦ RST : Connection Reset.
 - ♦ SYN: Sincronize Sequence Number.
 - ♦ FIN : Source closing connection.
- *Urgent Pointer* field defines the portion of data that should be considered urgent.
 - ♦ Not used by modern protocols.
- *Checksum* is used to detect errors.



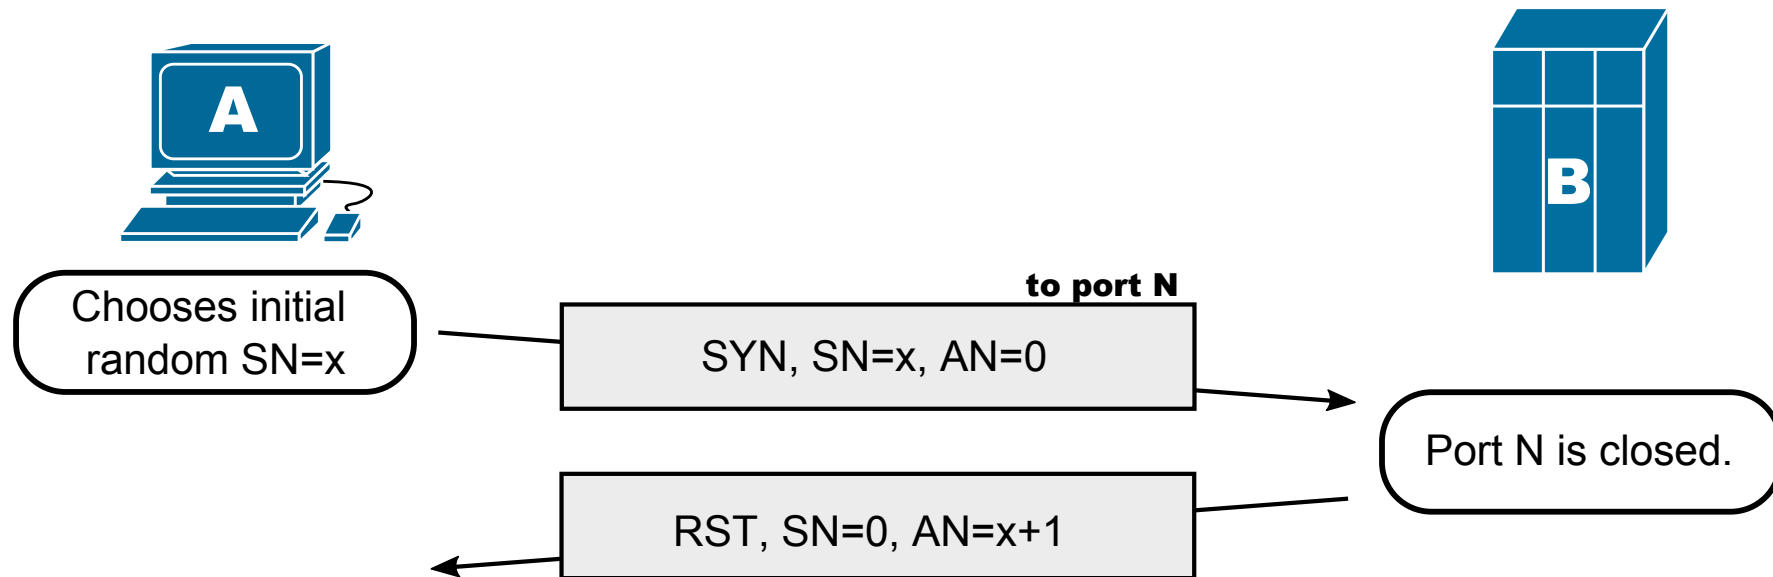
Establishment of a TCP Connection

- 3-Way Handshake.


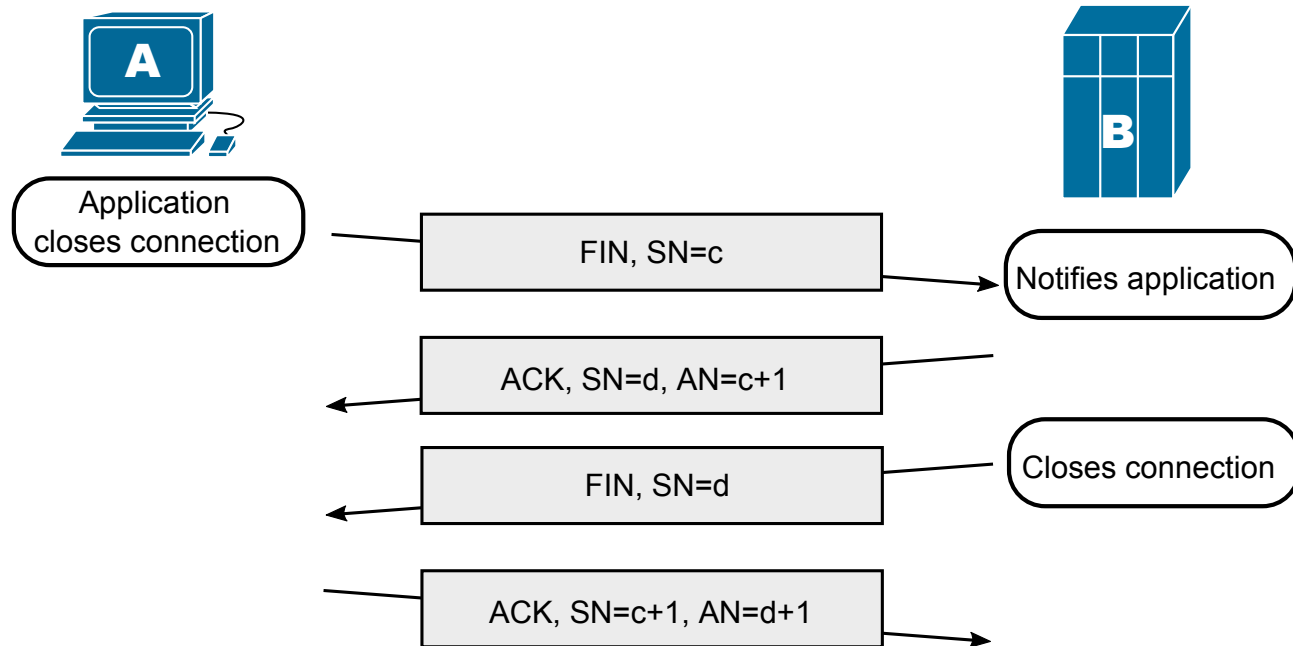
- Synchronizes the both end-points initial *Sequence Numbers* (SYN), and acknowledges it (ACK flag).



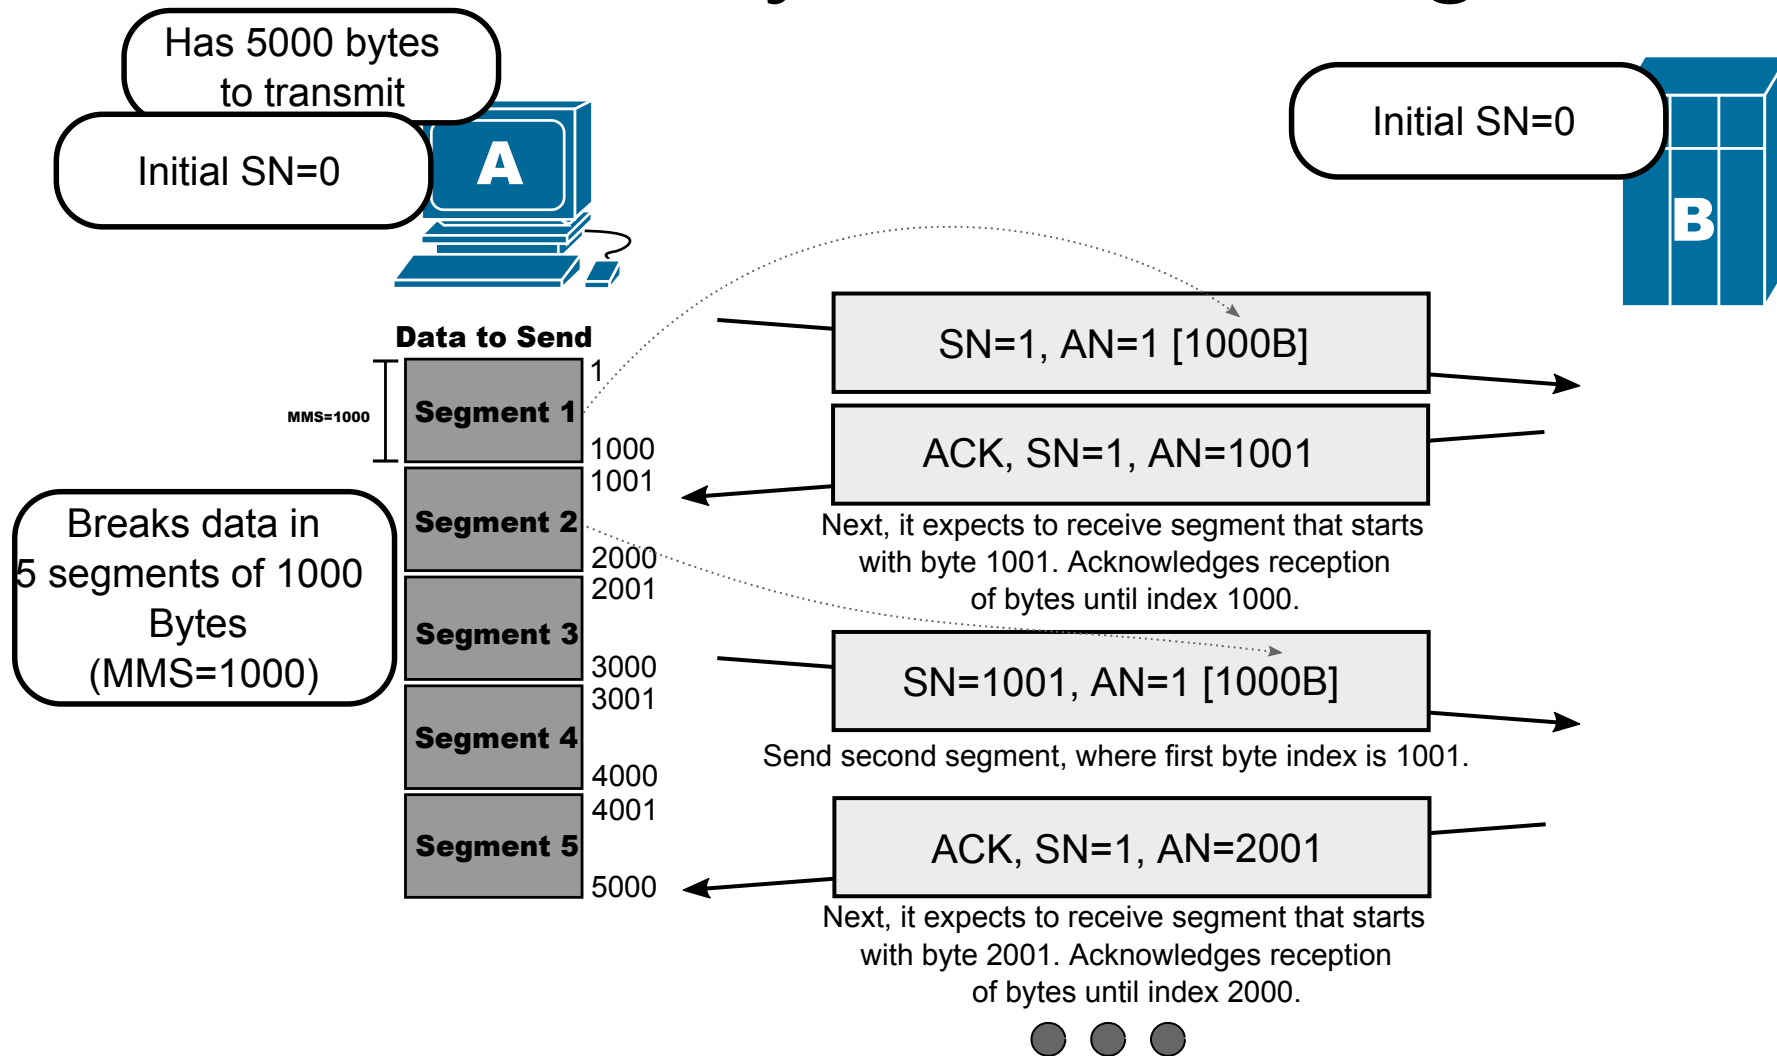
Establishment of a TCP Connection to a Closed Port



Closing a TCP Connection



Data Delivery Acknowledgment




- Sender can (usually) send more than one packet before receiving the ACK.
 - Depends on the congestion window (next slide).
- Both end-points can send data.
 - A packet with a data segment, can acknowledge the reception of data a segment received from the other end-point.

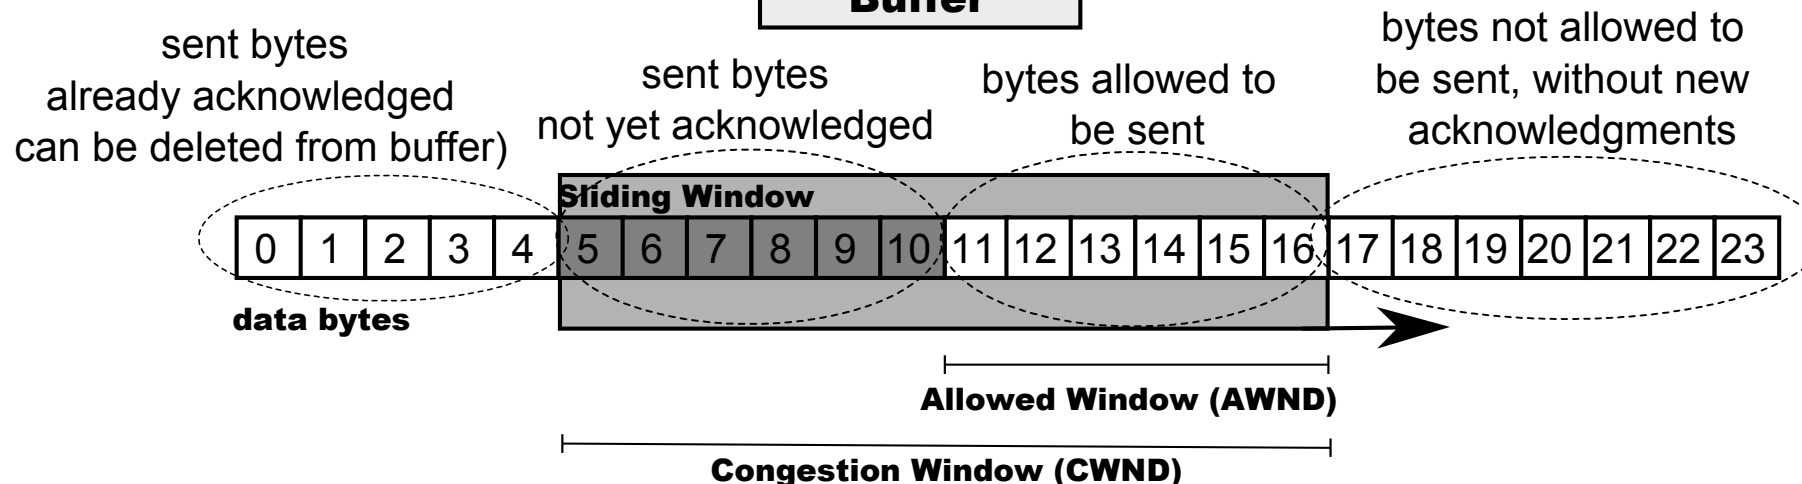


TCP Congestion Control (1)


- Uses a sliding window to determine the number of packets/bytes the sender is allowed to transmit.



- Congestion Window (CWND)
 - ◆ Set by sender to avoid overloading network.
 - ◆ The maximum value of CWND is RWND.
- Allowed Window (AWND)
 - ◆ Number of bytes allowed to be sent.
 - ◆ $AWND = \min(CWND - NACK, RWND)$
 - ➔ NACK: Not Acknowledged bytes



TCP Congestion Control (2)



Has 4000 bytes to transmit

Both have:
MSS=1000 bytes (by configuration)
RWND=2500 (Window value after establishment)
CWND=RWND



CWND=2500, NACK=0, AWND=2500

CWND=2500, NACK=1000, AWND=1500

CWND=2500, NACK=2000, AWND=500

CWND=2500, NACK=2500 AWND=0

CWND=2500, NACK=1500 AWND=0 (RWND=1500)

CWND=2500, NACK=500 AWND=0 (RWND=500)

CWND=2500, NACK=0 AWND=0 (RWND=0)

SN=x+1, AN=y+1 [1000B]

SN=x+1001, AN=y+1 [1000B]

SN=x+2001, AN=y+1 [500B]

ACK, SN=y+1, AN=y+1001, W=1500 [0B]

ACK, SN=y+1, AN=y+2001, W=500 [0B]

ACK, SN=y+1, AN=y+2501, W=0 [0B]

Transfers 2500B to application.

CWND=2500, NACK=0 AWND=2500 (RWND=2500)

CWND=2500, NACK=1000, AWND=1500

CWND=2500, NACK=1500, AWND=1000

SN=x+2501, AN=y+1 [1000B]

SN=x+3501, AN=y+1 [500B]

ACK, SN=y+1, AN=y+2501, W=2500 [0B]

ACK, SN=y+1, AN=y+3501, W=1500 [0B]

ACK, SN=y+1, AN=y+4001, W=500 [0B]




TCP Congestion Control (3)

- A packet with a data segment is considered lost when:
 - ♦ Timeout: After some time no ACK is received for that data segment
 - ♦ With Fast Retransmit/Recovery: After 3 or more duplicate ACKs are received for the previous data segment.
- When a packet is lost, TCP automatically decreases transmission rate to adapt to network conditions.
 - ♦ i.e., Decreases CWND size.
- When data delivery is acknowledged, TCP increases transmission rate.
 - ♦ i.e., Increases CWND size.
- The way the CWND size varies depend on the TCP Algorithms used:
 - ♦ Tahoe (Original TCP, 1988) uses:
 - Fast Retransmit, Slow Start, and Congestion Avoidance.
 - ♦ Reno (1990) uses:
 - Uses Fast Recovery, and Congestion Avoidance.
- At the beginning, the initial CWND value is usually 2, 3, 4, or 10 MSS and then the terminal starts the *Slow Start* with $SSThresh = RWND$.



TPC Fast Retransmit

- A segment is considered lost if 3 or more duplicate ACKs are received for the previous data segment.
 - ♦ Faster detection than waiting for a timeout.
 - ♦ Requires receiver to work.
- TCP retransmits immediately the lost segment.
- The TCP algorithm enters Slow-Start, with:
 - ♦ $SSThresh = CWND/2$
 - ♦ $CWND = Initial\ CWND$




TCP Slow-Start

- At the beginning, the initial CWND value is usually 2, 3, 4, or 10 MSS and the terminal starts the *Slow Start* with $SSThresh = RWND$.
- CWND size grows very fast while smaller than the predefined threshold ($CWND < SSThresh$).
 - ◆ When a ACK arrives the CWND is updated: $CWND = CWND + N$,
 - N is the number of bytes acknowledged in the ACK.
 - Results that the window size (approximately) doubles each round-trip time.
 - Exponential growth.
 - Continues until a loss occurs or CWND reaches RWND.
- When a loss occurs, $SSThresh$ is defined as $CWND/2$ and Slow-Start begins again from its initial CWND.
- Once the CWND reaches the $SSThresh$, it changes to congestion avoidance algorithm.
 - ◆ Linear growth.



TCP Congestion Avoidance


- When a ACK arrives the CWND is updated:
 $CWND = CWND + N/CWND$,
 - ◆ This results in a linear increase of the CWND.



- ◆ RWND is 64, and initial CWND is 4 MSS (segments).
- ◆ Initial SSThresh=RWND=64. After first loss: SSThresh=CWND/2=32. After second loss: SSThresh=CWND/2=26.

TPC Fast Recovery

- The same as TPC Fast Retransmit.
- However when a loss occurs enters directly to Congestion Avoidance with:
 - $SSThresh = CWND/2$
 - $CWND = SSThresh$
 - ➔ Some implementation have: $CWND = SSThresh + 3 * MSS$.




Other TCP Algorithms

- NewReno (1996)
 - ◆ Allows for partial ACK.
 - ◆ When a loss occurs, CWND is defined as $\beta \cdot \text{CWND}$, with $\beta=0.5$. When a ACK arrives, CWND is updated as $\text{CWND}=\text{CWND}+\alpha$, with $\alpha=1$ MSS.
 - ◆ Used by default in Windows and supported by Mac OS X.
 - Used in Windows XP and earlier.
 - After Windows Vista, Compound TCP can also be enabled.
- CUBIC (2005)
 - ◆ Uses a cubic function to control the CWND.
 - ◆ Used by Linux (kernel 2.6.19 and later) and supported by Mac OS X.
- Compound TCP (2006)
 - ◆ Adapts its behavior by use of a scalable delay-based component. T
 - Increases throughput more quickly in the congestion avoidance phase.
 - ◆ The AWND depend on the RTT measurements from successfully acknowledged packets.
 - ◆ Windows OS supports it as an option.
- Low Extra Delay Background Transport (LEDBAT)
 - ◆ Delay-based congestion control algorithm that uses all the available bandwidth while limiting the increase in delay. Measures one-way delay.
 - ◆ Supported by Windows 10 and latest versions of Mac OS X.



TCP Algorithms Comparison

- Tahoe




- Behavior and performance depends greatly on traffic and losses patterns, and end-points' behaviors


- In these examples:

- Using a traffic pattern that sends a number of packets equal to the CWND, in one second interval, equally spaced over the interval.
- Last packet of interval 8s-9s and 32s-33s are lost.
- ACWD is 64, and initial CWND is 4 MSS (segments).
- Receiver sends one ACK per packet.
- Each dot represents a received ACK.
- Tahoe sends 1585 packets, Reno 2017 packets and NewReno 2938 packets.


- Reno



- NewReno



- Cubic



Multipath TCP (MPTCP)

- TCP is essentially a single-path protocol.
 - When a TCP connection is established, the transmission is bound to the IP addresses of the two end-points.
 - If one address changes the TCP session will fail.
 - TCP can not load balance segments using more than um TCP session.
 - ➔ This load balancing must be done at the application level.
- Multipath TCP allows multiple subflows within a single MPTCP session.
 - A MPTCP session starts with an initial subflow, using the traditional 3-Way Handshake.
 - After the first MPTCP subflow is established, additional subflows can also established similar to the tradional TCP 3-Way Handshake. However, but rather than being a separate session, all subflows are bounded to the same MPTCP session.
 - Data can then be sent over any of the active subflows, using joint *Sequence* and *Acknowledgment Numbers*.
- Apple's Siri application uses Multipath TCP.

